

3_ex_regularized-regression-commented

June 27, 2023

1 3 Automatic feature selection for LDA as regression

1.1 3.1 Implement Orthogonal Matching Pursuit (8 points)

```
[1]: import copy
import numpy as np
from scipy.linalg import lstsq
```

Essentially the same solution. However we don't return A. We also preallocate A and B and return early if there is no residual left. We compute the residual and print it at every step. We also use scipy to solve the least squares problem. We compute the inactive betas explicitly.

```
[2]: def omp_regression(X, y, T):
    """Orthogonal Matching Pursuit is a simple greedy sparse regression
    algorithm. It approximates
    the exact algorithms for least squares under L1 regularization

    :param X: Input Matrix of shape  $R \times D$ 
    :param y: Output Vector of shape  $R \times 1$ 
    :param T: The desired number of non zero elements in the final solution
    :return beta_ts, optimal_t: The solution weights at each time t and the  $t_{\text{opt}}$ 
    where the solution is closest to the real solution.
    """

    assert T > 0, "Number of non zero Elements is smaller than 0"

    N = X.shape[0]
    D = X.shape[1]
    assert N == y.shape[0], "Dimension of Inputs and output does not match."

    # active matrix X
    X_active = np.zeros(X.shape)
    # inactive matrix
    X_inactive = copy.deepcopy(X)
    # zero matrix for generating inactive matrix
    zero_matrix = np.zeros(X.shape)
    # beta solution list
    beta_ts = np.zeros((D, T))
```

```

A = np.array(D*[False])
B = np.array(D*[True])
r = y
residual_norms = []

for t in range(T):
    # 1
    # Find most active column or
    # maximal correlation with the current residual
    correlation = np.abs(np.dot(X.T, r))
    j = np.argmax(correlation)

    # 2
    # set active index to 1
    A[j] = True
    # remove active index from B and set to 0
    B[j] = False

    # 3
    # Select active A
    X_active[:,A] = X[:,A]
    X_inactive[:,A] = zero_matrix[:,A]

    # 4
    # Calculate least squares
    beta_t, residue, rank, singular_value = lstsq(X_active, y)

    # 5
    # Update the residual
    r = y - np.dot(X_active, beta_t)
    residual_norms.append(np.linalg.norm(r))

    # Stop early if solution is found
    if np.sum(r) == 0:
        break

    beta_ts[:, t] = beta_t
    print(f"Distance to optimal solution on train after {t+1} steps:␣
↪{residual_norms[-1]}")
    print("#"*40)

return beta_ts

```

```

[3]: X = np.random.randint(5, size=(3,10))
y = np.random.randint(5, size=(3))
T = 3

```

```
[4]: solutions = omp_regression(X, y, T)
```

```
Distance to optimal solution on train after 1 steps: 1.1602387022306428
#####
Distance to optimal solution on train after 2 steps: 0.5015568278463086
#####
Distance to optimal solution on train after 3 steps: 1.2947314098277875e-15
#####
```

1.2 3.2 Classification with sparse LDA (8 points)

```
[5]: from sklearn.datasets import load_digits
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

We redistribute the data rather than balance it from the get go. Thus our dataset is complete compared to the sample solution.

```
[6]: def balance_training_split(X_train, X_test, y_train, y_test):
    """
    From a given train test split of binary data, balance training split.
    The split has to contain both labels for this algorithm to work.

    :param X_train: M x D, the training feature instances
    :param X_test: M x D, the test feature instances
    :param y_train: M x 1, the training labels
    :param y_test: M x 1, the test features
    """
    labels = np.unique(y_train)
    assert len(labels) == 2, "The split has to contain both labels for this_
algorithm to work."

    train_class_0_idx = np.where(y_train == labels[0])[0]
    train_class_1_idx = np.where(y_train == labels[1])[0]
    test_class_0_idx = np.where(y_test == labels[0])[0]
    test_class_1_idx = np.where(y_test == labels[1])[0]
    diff = np.abs(len(train_class_0_idx) - len(train_class_1_idx))

    to_test_idx = []
    to_train_idx = []
    if diff % 2 != 0:
        if len(train_class_0_idx) > len(train_class_1_idx):
            to_test_idx.append(train_class_0_idx[0])
            train_class_0_idx = np.delete(train_class_0_idx, 0)
        else:
            to_test_idx.append(train_class_1_idx[0])
            train_class_1_idx = np.delete(train_class_1_idx, 0)
```

```

diff -= 1

if diff > 0:
    switched_elements = list(range(int(diff/2)))
    if len(train_class_0_idx) > len(train_class_1_idx):
        to_test_idx += train_class_0_idx[switched_elements].tolist()
        train_class_0_idx = np.delete(train_class_0_idx, switched_elements)
        to_train_idx = test_class_1_idx[switched_elements].tolist()
        test_class_1_idx = np.delete(test_class_1_idx, switched_elements)
    else:
        to_test_idx += train_class_1_idx[switched_elements].tolist()
        train_class_1_idx = np.delete(train_class_1_idx, switched_elements)
        to_train_idx = test_class_0_idx[switched_elements].tolist()
        test_class_0_idx = np.delete(test_class_0_idx, switched_elements)

y_to_test = y_train[to_test_idx]
X_to_test = X_train[to_test_idx]
y_to_train = y_test[to_train_idx]
X_to_train = X_test[to_train_idx]

y_train = np.delete(y_train, to_test_idx)
X_train = np.delete(X_train, to_test_idx, axis=0)
y_test = np.delete(y_test, to_train_idx)
X_test = np.delete(X_test, to_train_idx, axis=0)

y_train = np.concatenate([y_train, y_to_train])
X_train = np.concatenate([X_train, X_to_train], axis=0)
y_test = np.concatenate([y_test, y_to_test])
X_test = np.concatenate([X_test, X_to_test], axis=0)

return X_train, X_test, y_train, y_test

```

```

[7]: # load and select data
digits = load_digits()

data = digits["data"]
images = digits["images"]
target = digits["target"]
target_names = digits["target_names"]

# filter data to images with label 3 or 9
idx = (target == 3) | (target == 9)
data_cleaned = data[idx]
images_cleaned = images[idx]

```

```

target_cleaned = target[idx]

# relabel labels 3 and 9 to -1, 1
idx_target_3 = (target_cleaned == 3)
idx_target_9 = (target_cleaned == 9)
target_cleaned[idx_target_3] = -1
target_cleaned[idx_target_9] = 1

# split the data
X_train, X_test, y_train, y_test = train_test_split(data_cleaned,
    ↪target_cleaned)

# balance training data
X_train, X_test, y_train, y_test = balance_training_split(X_train, X_test,
    ↪y_train, y_test)

```

The data standardization is basically the same with the exception that we fill NaN values with 0.

```

[8]: def standardize_data(data: np.ndarray):
    """
    Do data standardization:  $x_{\text{standardized}} = (x - \text{mean}(x)) / \text{std}(x)$ 

    :param data: N x D ndarray containing the data to be standarized.
    """
    centralized_data = data - np.mean(data, axis=0)
    std_dev = np.std(data, axis=0)
    standardized_data = centralized_data / std_dev
    # For dimensions with std_dev = 0 choose standardized_data = 0
    # For these entries the Z-score is defined as 0
    zero_std_dev = std_dev == 0
    standardized_data[:, zero_std_dev] = 0
    return standardized_data

```

```

[9]: data_cleaned_standardized = standardize_data(data_cleaned)

scaler = StandardScaler()
data_cleaned_sklearn_standardized = scaler.fit_transform(data_cleaned)

# verify if our solution is correct
print(f"Our standard scaling standard deviation and mean:\n{np.
    ↪std(data_cleaned_standardized)}, {np.mean(data_cleaned_standardized)}")
print(f"Sklearn's standard scaling standard deviation and mean:\n{np.
    ↪std(data_cleaned_sklearn_standardized)}, {np.
    ↪mean(data_cleaned_sklearn_standardized)}")

```

Our standard scaling standard deviation and mean:

0.9354143466934854, -1.5292328162880945e-19

Sklearn's standard scaling standard deviation and mean:

```
0.9354143466934854, -1.5292328162880945e-19
```

```
/tmp/ipykernel_17617/2403962991.py:9: RuntimeWarning: invalid value encountered  
in true_divide
```

```
    standardized_data = centralized_data / std_dev
```

```
[10]: # Scale the test and training data  
X_standardized_train = standardize_data(X_train)  
X_standardized_test = standardize_data(X_test)
```

```
/tmp/ipykernel_17617/2403962991.py:9: RuntimeWarning: invalid value encountered  
in true_divide
```

```
    standardized_data = centralized_data / std_dev
```

```
[11]: # for non-standardized images  
solutions_digits = omp_regression(X_train, y_train, X_train.shape[1])
```

```
Distance to optimal solution on train after 1 steps: 13.915025701524401  
#####  
Distance to optimal solution on train after 2 steps: 8.870110705921896  
#####  
Distance to optimal solution on train after 3 steps: 8.540165900419137  
#####  
Distance to optimal solution on train after 4 steps: 7.78826891685411  
#####  
Distance to optimal solution on train after 5 steps: 7.278822350683555  
#####  
Distance to optimal solution on train after 6 steps: 7.002903782491763  
#####  
Distance to optimal solution on train after 7 steps: 6.4968017600884105  
#####  
Distance to optimal solution on train after 8 steps: 6.263334102419298  
#####  
Distance to optimal solution on train after 9 steps: 6.12400397794607  
#####  
Distance to optimal solution on train after 10 steps: 5.969128638107378  
#####  
Distance to optimal solution on train after 11 steps: 5.787048695539513  
#####  
Distance to optimal solution on train after 12 steps: 5.647405689998397  
#####  
Distance to optimal solution on train after 13 steps: 5.552547880539769  
#####  
Distance to optimal solution on train after 14 steps: 5.488524793585701  
#####  
Distance to optimal solution on train after 15 steps: 5.385644913883237  
#####  
Distance to optimal solution on train after 16 steps: 5.2882567450780105  
#####
```

Distance to optimal solution on train after 17 steps: 5.232814856977235

Distance to optimal solution on train after 18 steps: 5.1550932048944516

Distance to optimal solution on train after 19 steps: 5.11441615513987

Distance to optimal solution on train after 20 steps: 5.0360981650891565

Distance to optimal solution on train after 21 steps: 4.99392270821011

Distance to optimal solution on train after 22 steps: 4.958227629304

Distance to optimal solution on train after 23 steps: 4.904588542992632

Distance to optimal solution on train after 24 steps: 4.890791791310587

Distance to optimal solution on train after 25 steps: 4.873410209340849

Distance to optimal solution on train after 26 steps: 4.844963600900127

Distance to optimal solution on train after 27 steps: 4.731437635099091

Distance to optimal solution on train after 28 steps: 4.703736258738175

Distance to optimal solution on train after 29 steps: 4.6676105825416485

Distance to optimal solution on train after 30 steps: 4.657558773804667

Distance to optimal solution on train after 31 steps: 4.653762182764904

Distance to optimal solution on train after 32 steps: 4.641261631402614

Distance to optimal solution on train after 33 steps: 4.627736432511158

Distance to optimal solution on train after 34 steps: 4.612911525714947

Distance to optimal solution on train after 35 steps: 4.60272483077835

Distance to optimal solution on train after 36 steps: 4.59984155259062

Distance to optimal solution on train after 37 steps: 4.59459907238186

Distance to optimal solution on train after 38 steps: 4.589251670975272

Distance to optimal solution on train after 39 steps: 4.58889902802033

Distance to optimal solution on train after 40 steps: 4.57761834864538
#####

Distance to optimal solution on train after 41 steps: 4.5728637276361335

Distance to optimal solution on train after 42 steps: 4.570672315828692

Distance to optimal solution on train after 43 steps: 4.566472907087987

Distance to optimal solution on train after 44 steps: 4.566261261848185

Distance to optimal solution on train after 45 steps: 4.562445536738479

Distance to optimal solution on train after 46 steps: 4.561209584873174

Distance to optimal solution on train after 47 steps: 4.560878101173517

Distance to optimal solution on train after 48 steps: 4.560733083907693

Distance to optimal solution on train after 49 steps: 4.556604320867482

Distance to optimal solution on train after 50 steps: 4.551765630197875

Distance to optimal solution on train after 51 steps: 4.542792822192767

Distance to optimal solution on train after 52 steps: 4.542666652769874

Distance to optimal solution on train after 53 steps: 4.5426350748404865

Distance to optimal solution on train after 54 steps: 4.542630982972832

Distance to optimal solution on train after 55 steps: 4.542630982972832

Distance to optimal solution on train after 56 steps: 4.542630982972832

Distance to optimal solution on train after 57 steps: 4.542630982972832

Distance to optimal solution on train after 58 steps: 4.542630982972832

Distance to optimal solution on train after 59 steps: 4.542630982972832

Distance to optimal solution on train after 60 steps: 4.542630982972832

Distance to optimal solution on train after 61 steps: 4.542630982972832

Distance to optimal solution on train after 62 steps: 4.542630982972832

Distance to optimal solution on train after 63 steps: 4.542630982972832

Distance to optimal solution on train after 64 steps: 4.542630982972832
#####


```
[12]: # for standardized images
      solutions_standardized_digits = omp_regression(X_standardized_train, y_train,
      ↪X_standardized_train.shape[1])
```

```
Distance to optimal solution on train after 1 steps: 10.881641649376581
#####
Distance to optimal solution on train after 2 steps: 8.854083917667381
#####
Distance to optimal solution on train after 3 steps: 7.856843624497157
#####
Distance to optimal solution on train after 4 steps: 7.338428098819996
#####
Distance to optimal solution on train after 5 steps: 6.828241236680242
#####
Distance to optimal solution on train after 6 steps: 6.32592592713819
#####
Distance to optimal solution on train after 7 steps: 6.16284441377525
#####
Distance to optimal solution on train after 8 steps: 5.9979524347077104
#####
Distance to optimal solution on train after 9 steps: 5.897293046572178
#####
Distance to optimal solution on train after 10 steps: 5.685591858717865
#####
Distance to optimal solution on train after 11 steps: 5.597920668477057
#####
Distance to optimal solution on train after 12 steps: 5.4794353322083635
#####
Distance to optimal solution on train after 13 steps: 5.382611982858704
#####
Distance to optimal solution on train after 14 steps: 5.261317280688126
#####
Distance to optimal solution on train after 15 steps: 5.210579188393709
#####
Distance to optimal solution on train after 16 steps: 5.145102336830046
#####
Distance to optimal solution on train after 17 steps: 5.065430893356197
#####
Distance to optimal solution on train after 18 steps: 5.023569972568772
#####
Distance to optimal solution on train after 19 steps: 4.980650998596544
#####
Distance to optimal solution on train after 20 steps: 4.915051077477012
#####
Distance to optimal solution on train after 21 steps: 4.826260971175179
#####
Distance to optimal solution on train after 22 steps: 4.769694454372308
#####
```

Distance to optimal solution on train after 23 steps: 4.73078877023645

Distance to optimal solution on train after 24 steps: 4.708621032120979

Distance to optimal solution on train after 25 steps: 4.680238620073461

Distance to optimal solution on train after 26 steps: 4.656316678505915

Distance to optimal solution on train after 27 steps: 4.648475013990084

Distance to optimal solution on train after 28 steps: 4.6427648908548

Distance to optimal solution on train after 29 steps: 4.636121621717053

Distance to optimal solution on train after 30 steps: 4.6251783550768435

Distance to optimal solution on train after 31 steps: 4.6164504833679185

Distance to optimal solution on train after 32 steps: 4.604081184697056

Distance to optimal solution on train after 33 steps: 4.592523984487908

Distance to optimal solution on train after 34 steps: 4.587910092535218

Distance to optimal solution on train after 35 steps: 4.5719292908842935

Distance to optimal solution on train after 36 steps: 4.566867782449537

Distance to optimal solution on train after 37 steps: 4.558725211229966

Distance to optimal solution on train after 38 steps: 4.554547762484517

Distance to optimal solution on train after 39 steps: 4.5384370717818365

Distance to optimal solution on train after 40 steps: 4.534983873427219

Distance to optimal solution on train after 41 steps: 4.531974965873104

Distance to optimal solution on train after 42 steps: 4.528527368837843

Distance to optimal solution on train after 43 steps: 4.527532655321474

Distance to optimal solution on train after 44 steps: 4.526831740680015

Distance to optimal solution on train after 45 steps: 4.525647142647113

Distance to optimal solution on train after 46 steps: 4.5250904141956285
#####

```

Distance to optimal solution on train after 47 steps: 4.524540391313629
#####
Distance to optimal solution on train after 48 steps: 4.5243221637308535
#####
Distance to optimal solution on train after 49 steps: 4.524213414000305
#####
Distance to optimal solution on train after 50 steps: 4.649242084123436
#####
Distance to optimal solution on train after 51 steps: 4.649242084123436
#####
Distance to optimal solution on train after 52 steps: 4.649242084123436
#####
Distance to optimal solution on train after 53 steps: 4.649242084123436
#####
Distance to optimal solution on train after 54 steps: 4.649242084123436
#####
Distance to optimal solution on train after 55 steps: 4.649242084123436
#####
Distance to optimal solution on train after 56 steps: 4.649242084123436
#####
Distance to optimal solution on train after 57 steps: 4.649242084123436
#####
Distance to optimal solution on train after 58 steps: 4.649242084123436
#####
Distance to optimal solution on train after 59 steps: 4.649242084123436
#####
Distance to optimal solution on train after 60 steps: 4.649242084123436
#####
Distance to optimal solution on train after 61 steps: 4.649242084123436
#####
Distance to optimal solution on train after 62 steps: 4.649242084123436
#####
Distance to optimal solution on train after 63 steps: 4.649242084123436
#####
Distance to optimal solution on train after 64 steps: 4.649242084123436
#####

```

```

[13]: def predict(solutions, X_train, X_test):
        # we can assume that the mean over all is the same as the mean over each
        ↪class
        # since there are balanced in train
        mu = np.mean(X_train, axis=0)
        lhs = np.dot(X_test, solutions)
        rhs = np.dot(mu, solutions)
        return np.sign(lhs - rhs)

def error_rates(predictions, y_test):

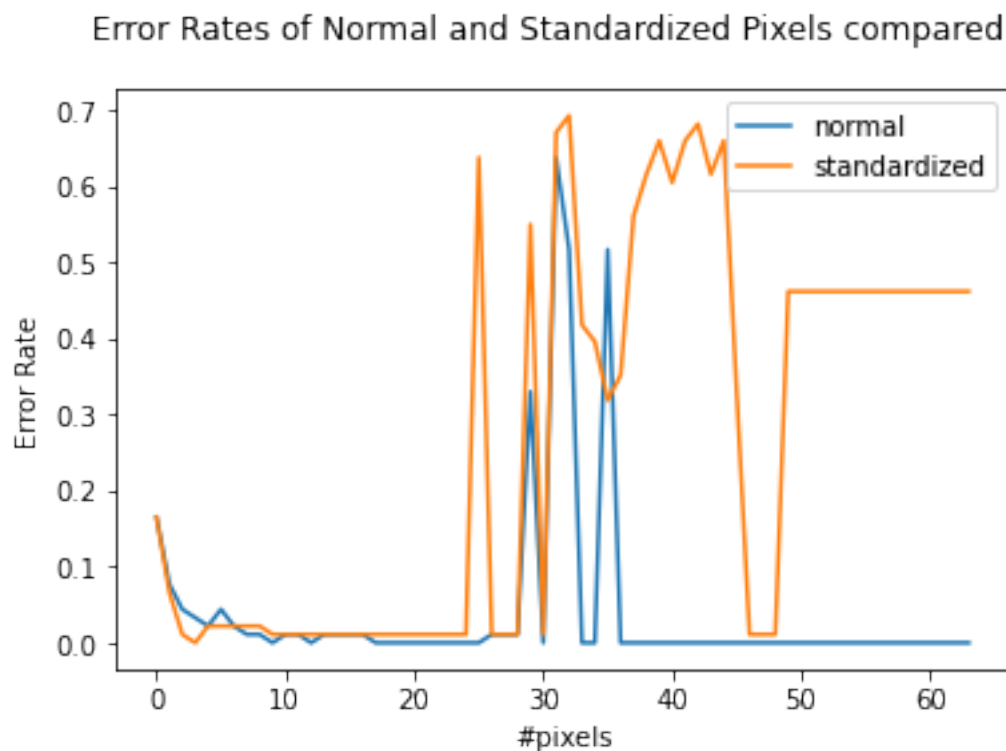
```

```
errors = np.abs((predictions.T - y_test)/2)
return np.mean(errors, axis=1)
```

```
[14]: errors = error_rates(predict(solutions_digits, X_train, X_test), y_test)
error_rates_standardized = error_rates(predict(solutions_standardized_digits,
↪X_standardized_train, X_standardized_test), y_test)
```

Our prediction method is wrong therefore our error rates look very different to sample solution. Activation order missing.

```
[15]: plt.plot(list(range(len(errors))), errors, label="normal")
plt.plot(list(range(len(error_rates_standardized))), error_rates_standardized,
↪label="standardized")
plt.legend()
plt.suptitle("Error Rates of Normal and Standardized Pixels compared")
plt.xlabel("#pixels")
plt.ylabel("Error Rate")
plt.show()
```



Between 10 and 20 pixels we have a high probability to get a perfect classifier. With 3 to 5 pixels we see an error rate of at or below 5 percent.

Standardization did not give us significant improvements in classification quality though we suspect that it could be more stable in some scenarios.

Unfortunately, we were not able to finish the pixel visualization in time. We would propose an animated image which activates the pixels according if the respective dimensions in `beta_t` are nonzero. Then we would look at an averaged 3 image and decide whether the respective beta votes for or against 3 by looking if the average brightness is smaller or larger than some threshold (e.g. 0.2 of the max brightness).

Depending on how well we chose the pixels in exercise 1 we would see overlap with the solution here since OMP will choose the most optimal pixels for the classification task first.